

Objective: Develop a Classifier to predict whether a customer is going to open a deposit account.

Data: Bank Dataset (bank-test.csv; bank-train.csv)

See UCI Repository <https://archive.ics.uci.edu/ml/datasets/Bank+Marketing> (**with variable 11 – duration, excluded**) for more information

Part1: Build four separate classifiers, and compare their performance.

```
library(caret)
library(plyr)
library(dplyr)
library(C50)
library(kernlab)
library(RWeka)
trainingData <- read.csv("C:\\Users\\ludai\\Desktop\\BAN620\\Assignment3\\bank-training.csv", row.names = 1)
testData <- read.csv("C:\\Users\\ludai\\Desktop\\BAN620\\Assignment3\\bank-test.csv", row.names = 1)
nrow(trainingData)
nrow(testData)
prop.table(table(trainingData$y))
prop.table(table(testData$y))

[1] 3090
> nrow(testData)
[1] 1029
> prop.table(table(trainingData$y))

      no      yes
0.8902913 0.1097087
> prop.table(table(testData$y))

      no      yes
0.8911565 0.1088435
```

a) Decision Tree

```
TrainingParameters <- trainControl(method = "cv", number = 10, repeats = 5)
DecTreeModel <- caret::train(trainingData[, -20], trainingData$y,
                             method = "C5.0",
                             trControl = TrainingParameters,
                             na.action = na.omit
                             )
DecTreeModel
summary(DecTreeModel)
#Testing the Model
DTPredictions <- predict(DecTreeModel, testData, na.action = na.pass)
confusionMatrix(DTPredictions, testData$y)
confusionMatrix
```

```
> DecTreeModel
C5.0

3090 samples
 19 predictor
 2 classes: 'no', 'yes'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 2781, 2781, 2781, 2782, 2781, 2781, ...
Resampling results across tuning parameters:

model  winnow  trials  Accuracy  Kappa
rules  FALSE   1       0.9019373 0.2603314
rules  FALSE  10       0.8970840 0.2761122
rules  FALSE  20       0.8980570 0.2905335
rules  TRUE    1       0.8996720 0.2586340
rules  TRUE   10       0.8974045 0.2706525
rules  TRUE  20       0.8993494 0.3038061
tree   FALSE   1       0.8996720 0.2647905
tree   FALSE  10       0.8967541 0.2812278
tree   FALSE  20       0.8980486 0.2909091
tree   TRUE    1       0.8996720 0.2644463
tree   TRUE   10       0.9003108 0.3024265
tree   TRUE  20       0.8983712 0.3050800

Accuracy was used to select the optimal model using the largest value.
The final values used for the model were trials = 1, model = rules and winnow = FALSE.
```

```

Confusion Matrix and Statistics

          Reference
Prediction no yes
no      911  98
yes       6  14

      Accuracy : 0.8989
      95% CI   : (0.8789, 0.9167)
No Information Rate : 0.8912
P-Value [Acc > NIR] : 0.2281

      Kappa : 0.1852
McNemar's Test P-Value : <2e-16

      Sensitivity : 0.9935
      Specificity : 0.1250
      Pos Pred Value : 0.9029
      Neg Pred Value : 0.7000
      Prevalence : 0.8912
      Detection Rate : 0.8853
      Detection Prevalence : 0.9806
      Balanced Accuracy : 0.5592

      'Positive' Class : no

```

b) Naive Bayes

```

set.seed(100)
NBModel <- train(trainingData[,-20], trainingData$y, method = "nb", trControl = trainControl(method = "cv",
number = 10, repeats = 5))
NBModel
NBPredictions <- predict(NBModel, testData)
confusionMatrix(NBPredictions, testData$y)
> NBModel
Naive Bayes

3090 samples
19 predictor
2 classes: 'no', 'yes'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 2782, 2781, 2781, 2781, 2781, 2781, ...
Resampling results across tuning parameters:

usekernel Accuracy Kappa
FALSE      0.8424239 0.3291233
TRUE       0.8747759 0.3491600

Tuning parameter 'fL' was held constant at a value of 0
Tuning parameter 'adjust' was held constant at a value of 1
Accuracy was used to select the optimal model using the largest value.
The final values used for the model were fL = 0, usekernel = TRUE and adjust = 1.

```

```
> confusionMatrix(NBPredictions, testData$y)
Confusion Matrix and Statistics
```

```

      Reference
Prediction no yes
no      871  71
yes     46   41

      Accuracy : 0.8863
      95% CI   : (0.8653, 0.9051)
No Information Rate : 0.8912
P-Value [Acc > NIR] : 0.7121

      Kappa : 0.3502
McNemar's Test P-Value : 0.0265

      Sensitivity : 0.9498
      Specificity : 0.3661
      Pos Pred Value : 0.9246
      Neg Pred Value : 0.4713
      Prevalence : 0.8912
      Detection Rate : 0.8465
      Detection Prevalence : 0.9155
      Balanced Accuracy : 0.6580

      'Positive' Class : no
```

c) Support Vector Machines (SVM)

```
set.seed(120)
svm_model <- train(y~., data = trainingData,
  method = "svmPoly",
  trControl= trainControl(method = "cv", number = 10, repeats = 5),
  tuneGrid = data.frame(degree = 1, scale = 1, C = 1))
svm_model
SVMPredictions <- predict(svm_model, testData, na.action = na.pass)
confusionMatrix(SVMPredictions, testData$y)
> svm_model
Support Vector Machines with Polynomial Kernel

3090 samples
19 predictor
2 classes: 'no', 'yes'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 2781, 2780, 2781, 2781, 2782, 2781, ...
Resampling results:

      Accuracy   Kappa
0.8986969  0.2777773

Tuning parameter 'degree' was held constant at a value of 1
Tuning parameter 'scale' was held constant at a value of
1
Tuning parameter 'C' was held constant at a value of 1
```

```
> confusionMatrix(SVMPredictions, testData$y)
Confusion Matrix and Statistics
```

```

      Reference
Prediction no yes
no      906  98
yes     11  14

      Accuracy : 0.8941
      95% CI : (0.8736, 0.9122)
No Information Rate : 0.8912
P-Value [Acc > NIR] : 0.406

      Kappa : 0.1715
McNemar's Test P-Value : <2e-16

      Sensitivity : 0.9880
      Specificity : 0.1250
      Pos Pred Value : 0.9024
      Neg Pred Value : 0.5600
      Prevalence : 0.8912
      Detection Rate : 0.8805
      Detection Prevalence : 0.9757
      Balanced Accuracy : 0.5565

      'Positive' Class : no
```

d) Neural Network

```
nnmodel <- train(trainingData[,-20], trainingData$y, method = "nnet",
  trControl= trainControl(method = "cv", number = 10, repeats = 5))
nnmodel
nnetpredictions <- predict(nnmodel, testData, na.action = na.pass)
confusionMatrix(nnetpredictions, testData$y)
```

```
> nnmodel
Neural Network

3090 samples
19 predictor
2 classes: 'no', 'yes'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 2781, 2781, 2782, 2781, 2781, 2780, ...
Resampling results across tuning parameters:

  size  decay  Accuracy  Kappa
1      0e+00  0.8912598  0.0244378
1      1e-04  0.8902921  0.0000000
1      1e-01  0.9006481  0.3061129
3      0e+00  0.8948197  0.1305268
3      1e-04  0.8951465  0.0806496
3      1e-01  0.9003203  0.3125602
5      0e+00  0.8948228  0.0954000
5      1e-04  0.8974118  0.1148964
5      1e-01  0.8980559  0.3173886

Accuracy was used to select the optimal model using the largest value.
The final values used for the model were size = 1 and decay = 0.1.
> |
```

```
> confusionMatrix(nnetpredictions, testData$y)
Confusion Matrix and Statistics
```

```

      Reference
Prediction no yes
no      906  96
yes     11  16

      Accuracy : 0.896
      95% CI   : (0.8757, 0.914)
No Information Rate : 0.8912
P-Value [Acc > NIR] : 0.33

      Kappa : 0.1962
McNemar's Test P-Value : 4.639e-16

      Sensitivity : 0.9880
      Specificity : 0.1429
      Pos Pred Value : 0.9042
      Neg Pred Value : 0.5926
      Prevalence : 0.8912
      Detection Rate : 0.8805
      Detection Prevalence : 0.9738
      Balanced Accuracy : 0.5654

      'Positive' Class : no
```

Summary: Evaluation Four Model

```
> t <- matrix(cbind(c(0.8989,0.9935,0.1250,0.9029),c(0.8863,0.9498,0.3661,0.9246),
+ c(0.8941,0.9880,0.1250,0.9024),c(0.8960,0.9880,0.1429,0.9042)),
+ nrow=4,dimnames=list(c('Accuracy','Sensitivity/Recall','Specificity','Precision'),
+ c('DecTree','NB','SVM','NN')))
> t
```

	DecTree	NB	SVM	NN
Accuracy	0.8989	0.8863	0.8941	0.8960
Sensitivity/Recall	0.9935	0.9498	0.9880	0.9880
Specificity	0.1250	0.3661	0.1250	0.1429
Precision	0.9029	0.9246	0.9024	0.9042

```
> |
```

Based on the above matrix, we found the Decision Tree Model has the highest accuracy and recall. But in this case, we hope the recall(sensitivity) is low since the positive class is 'No', which means a customer does not want to open a deposit account. The Naive Bayes Model has the lowest recall. Therefore, it is hard to say which model is better only from the above measures.

Part2. F Measure weight justification, F Measures, Best Model

```
> model = c("DecTree","NB","SVM","NN")
> recall = c(0.9935,0.9498,0.9880,0.9880)
> precision = c(0.9029,0.9246,0.9024,0.9042)
> fscore <- 2 * precision * recall / (precision + recall)
> fscore_table = data.frame(model,recall,precision,fscore)
> fscore_table
  model recall precision  fscore
1 DecTree 0.9935    0.9029 0.9460358
2      NB 0.9498    0.9246 0.9370306
3      SVM 0.9880    0.9024 0.9432620
4       NN 0.9880    0.9042 0.9442444
> beta=as.numeric(0.5)
> x=beta^2
> x
[1] 0.25
> fmeasure=((1+x)*precision*recall)/(x*precision+recall)
> fmeasure_table=data.frame(model,recall,precision,fmeasure)
> fmeasure_table
  model recall precision  fmeasure
1 DecTree 0.9935    0.9029 0.9196735
2      NB 0.9498    0.9246 0.9295325
3      SVM 0.9880    0.9024 0.9183125
4       NN 0.9880    0.9042 0.9198031
> |

> beta=as.numeric(2)
> x=beta^2
> fmeasure=((1+x)*precision*recall)/(x*precision+recall)
> fmeasure_table=data.frame(model,recall,precision,fmeasure)
> fmeasure_table
  model recall precision  fmeasure
1 DecTree 0.9935    0.9029 0.9739540
2      NB 0.9498    0.9246 0.9446507
3      SVM 0.9880    0.9024 0.9696050
4       NN 0.9880    0.9042 0.9700200
. |
```

Precision can be thought of as a measure of exactness, whereas recall is a measure of completeness. In this case, the classification goal is to predict whether customers will subscribe to open deposit accounts. Here the positive class is 'no'. This means a customer does not subscribe to a deposit. So, it is important to choose a model with a low recall. In order to illustrate recall and precision for each model, we need to compute F-measure. The F measure is the harmonic mean of precision and recall. It gives equal weight to precision and recall. The F_β measure is a weighted measure of precision and recall. It assigns β times as much weight to recall as to precision. I assigned $\beta = 0.5$ (which weights precision twice as much as recall), and $\beta = 2$ (which weights recall twice as much as precision) respectively. Based on the analysis for the above table, Naive Bayes method is the recommended classification

method as it contains lowest recall and F score and the highest precision in the F measure table and F_2 measure table. We do not hope a model that aggressively classifies a customer response as 'no', we want more customers to open a bank account.

Part3. Cost Sensitive Model

```
library(C50)
cmatrix <- cbind(c(0,1), c(10,0))
cmatrix
trainingData <- read.csv("C:\\Users\\ludai\\Desktop\\BAN620\\Assignment3\\bank-training.csv",
row.names = 1)
testData<-read.csv("C:\\Users\\ludai\\Desktop\\BAN620\\Assignment3\\bank-test.csv", row.names = 1)
CostDTModel <- C5.0(y ~., data=trainingData, cost=cmatrix)
CostDTModel
summary(CostDTModel)
PredictionCostModel <- predict(CostDTModel, testData, na.action = na.pass)
confusionMatrix(PredictionCostModel, testData$y)
> CostDTModel
```

```
Call:
C5.0.formula(formula = y ~ ., data = trainingData, cost = cmatrix)
```

```
Classification Tree
Number of samples: 3090
Number of predictors: 19
```

```
Tree size: 61
```

```
Non-standard options: attempt to group attributes
```

```
Cost Matrix:
```

```
      no yes
no    0  10
yes   1   0
> |
```

```
Confusion Matrix and Statistics
```

```
      Reference
Prediction no yes
no        605  39
yes       312  73
```

```
      Accuracy : 0.6589
      95% CI   : (0.629, 0.6879)
No Information Rate : 0.8912
P-Value [Acc > NIR] : 1
```

```
      Kappa : 0.1505
McNemar's Test P-Value : <2e-16
```

```
      Sensitivity : 0.6598
      Specificity : 0.6518
      Pos Pred Value : 0.9394
      Neg Pred Value : 0.1896
      Prevalence : 0.8912
      Detection Rate : 0.5879
      Detection Prevalence : 0.6259
      Balanced Accuracy : 0.6558
```

```
'Positive' Class : no
```

Profit=73*10(revenue)-(312+73) *1=345

Based on cost matrix and confusion matrix, since the cost of a phone call is 1 unit and lost business is 10units, the revenue will be 73*10(True Negative), and the cost is 1*(312+73). The profit is 345.

Part4: Extra credit

a) Models Correlation

```
econtrol <- trainControl(method="cv", number=5, summaryFunction = twoClassSummary, savePredictions=TRUE,
classProbs=TRUE)
models <- caretList(y ~., data=trainingData,
methodList=c("svmPoly", "nnet", "C5.0", "nb"),
trControl = econtrol )
results <- resamples(models)
results$values
results$metrics
summary(results)
mcr <- modelCor(results)
mcr
```

```
> summary(results)
```

```
Call:
summary.resamples(object = results)
```

```
Models: svmPoly, nnet, C5.0, nb
Number of resamples: 5
```

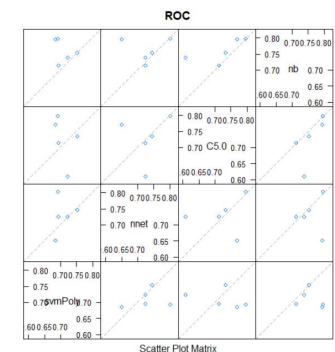
ROC	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
svmPoly	0.6849932	0.6912834	0.6942724	0.7091847	0.7228342	0.7525401	0
nnet	0.6497151	0.7238770	0.7242447	0.7288615	0.7450267	0.8014439	0
C5.0	0.6094652	0.7130218	0.7333155	0.7246515	0.7705699	0.7968850	0
nb	0.7129017	0.7381684	0.7523663	0.7589476	0.7939620	0.7973396	0

Sens	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
svmPoly	0.9818512	0.9836364	0.9836364	0.9854611	0.9872727	0.9909091	0
nnet	0.9564428	0.9745455	0.9836364	0.9771067	0.9836364	0.9872727	0
C5.0	0.9509091	0.9818182	0.9836364	0.9792767	0.9891107	0.9909091	0
nb	1.0000000	1.0000000	1.0000000	1.0000000	1.0000000	1.0000000	0

Spec	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
svmPoly	0.1323529	0.2352941	0.2352941	0.2273047	0.2500000	0.2835821	0
nnet	0.1470588	0.1764706	0.2500000	0.2509219	0.3134328	0.3676471	0
C5.0	0.1323529	0.2058824	0.2089552	0.2359087	0.2352941	0.3970588	0
nb	0.0000000	0.0000000	0.0000000	0.0000000	0.0000000	0.0000000	0

```
> mcr <- modelCor(results)
> mcr
```

	svmPoly	nnet	C5.0	nb
svmPoly	1.0000000	0.21416392	-0.4204132	-0.34446113
nnet	0.2141639	1.00000000	0.1514213	0.02842688
C5.0	-0.4204132	0.15142135	1.00000000	0.69102423
nb	-0.3444611	0.02842688	0.6910242	1.00000000



Based on result of Models correlation, I find these four models have low correlation (<0.75) by using resample and modelCor functions. C5.0 and NB have the highest positive correlation; and nnet and NB have the least correlation.

b) Ensemble Methods

```
# subset models to use
smallmodels <- c(models$svmPoly, models$C5.0, models$nb, models$nnet)
# Create stacked models. Try rpart
enstackmodel <- caretStack(models, method = "rpart")
print(enstackmodel)
enstackmodel <- caretStack(models, method = "C5.0", metric="Sens", trControl = trainControl(number = 5,
summaryFunction = twoClassSummary, classProbs = TRUE))
print(enstackmodel)
#Predict
enstackpredictions <- predict(enstackmodel, testData, na.action = na.omit)
# Create confusion matrix
cmBank <- confusionMatrix(enstackpredictions, testData$y, mode="everything", positive = "no")
cmBank
> print(enstackmodel)
A rpart ensemble of 2 base models: svmPoly, nnet, C5.0, nb
```

Ensemble results:
CART

3090 samples
4 predictor
2 classes: 'no', 'yes'

No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 3090, 3090, 3090, 3090, 3090, ..
Resampling results across tuning parameters:

cp	Accuracy	Kappa
0.006637168	0.8873470	0.2709128
0.008849558	0.8900828	0.2786850
0.058997050	0.8916057	0.1780602

```
> print(enstackmodel)
A C5.0 ensemble of 2 base models: svmPoly, nnet, C5.0, nb
```

Ensemble results:
C5.0

3090 samples
4 predictor
2 classes: 'no', 'yes'

No pre-processing
Resampling: Bootstrapped (5 reps)
Summary of sample sizes: 3090, 3090, 3090, 3090, 3090
Resampling results across tuning parameters:

model	winnow	trials	ROC	Sens	Spec
rules	FALSE	1	0.6190529	0.9845063	0.1830710
rules	FALSE	10	0.7421883	0.9751225	0.2272290
rules	FALSE	20	0.7471149	0.9753644	0.2486935
rules	TRUE	1	0.6191909	0.9787190	0.2504490
rules	TRUE	10	0.6646853	0.9752259	0.2595383
rules	TRUE	20	0.6646853	0.9752259	0.2595383
tree	FALSE	1	0.7034884	0.9757112	0.2082679
tree	FALSE	10	0.7436723	0.9728238	0.2472836
tree	FALSE	20	0.7473011	0.9724399	0.2438993
tree	TRUE	1	0.6641840	0.9775178	0.2519994
tree	TRUE	10	0.6906682	0.9771411	0.2489459
tree	TRUE	20	0.6906682	0.9771411	0.2489459

```
> cmBank
Confusion Matrix and Statistics
```

	Reference	
Prediction	no	yes
no	906	98
yes	11	14

Accuracy : 0.8941
95% CI : (0.8736, 0.9122)
No Information Rate : 0.8912
P-Value [Acc > NIR] : 0.406

Kappa : 0.1715
McNemar's Test P-Value : <2e-16

Sensitivity : 0.9880
Specificity : 0.1250
Pos Pred Value : 0.9024
Neg Pred Value : 0.5600
Precision : 0.9024
Recall : 0.9880
F1 : 0.9433
Prevalence : 0.8912
Detection Rate : 0.8805
Detection Prevalence : 0.9757
Balanced Accuracy : 0.5565

'Positive' Class : no

```
econtrol <- trainControl(method="cv", number=5, summaryFunction = twoClassSummary, savePredictions=TRUE,
classProbs=TRUE)
```

```
models <- caretList(y ~., data=trainingData,
methodList=c("nnet", "nb"),
trControl = econtrol)
smallmodels <- c(models$nb, models$nnet)
enstackmodel <- caretStack(models, method = "rpart")
```

```

print(enstackmodel)
enstackmodel <- caretStack(models, method = "C5.0", metric="Sens",trControl = trainControl(number = 5, summaryFunction =
twoClassSummary,classProbs = TRUE))
print(enstackmodel)
enstackpredictions <-predict(enstackmodel, testData, na.action = na.omit)
cmBank <-confusionMatrix(enstackpredictions, testData$y, mode="everything", positive = "no")
cmBank
A C5.0 ensemble of 2 base models: nnet, nb

Ensemble results:
C5.0

3090 samples
  2 predictor
  2 classes: 'no', 'yes'

No pre-processing
Resampling: Bootstrapped (5 reps)
Summary of sample sizes: 3090, 3090, 3090, 3090, 3090
Resampling results across tuning parameters:

```

Confusion Matrix and Statistics

```

      Reference
Prediction no yes
no      905  95
yes      12  17

      Accuracy : 0.896
      95% CI : (0.8757, 0.914)
      No Information Rate : 0.8912
      P-Value [Acc > NIR] : 0.33

      Kappa : 0.2056
      Mcnemar's Test P-Value : 2.241e-15

      Sensitivity : 0.9869
      Specificity : 0.1518
      Pos Pred Value : 0.9050
      Neg Pred Value : 0.5862
      Precision : 0.9050
      Recall : 0.9869
      F1 : 0.9442
      Prevalence : 0.8912
      Detection Rate : 0.8795
      Detection Prevalence : 0.9718
      Balanced Accuracy : 0.5693

      'Positive' Class : no

```

Ensemble model by using all four models:

Accuracy: 0.8941

Sensitivity: 0.9880

Specificity: 0.1250

Precision: 0.9024

Ensemble model by using least correlation nnet and nb models:

Accuracy: 0.8960

Sensitivity: 0.9869

Specificity: 0.1518

Precision: 0.9050

In ensemble method, we can combine predictors to improve classification accuracy. In this case, I used model stacking to create ensemble model. I find the accuracy is improved by using least correlation models.

C) Develop an ensemble model that returns a profit greater than 425 on the test data

```
> # Create models with sensitivity as optimization metric instead of accuracy
> ensmodel2 <- caretEnsemble(models,
+                             metric = "Sens",
+                             trControl = trainControl(number = 2, summaryFunction = twoClassSummary, classProbs = TRUE)
+ )
> summary(ensmodel2)
The following models were ensemble: svmPoly, nnet, C5.0, nb
They were weighted:
-2.9794 -0.2833 3.7614 2.4072 -24190.0965
The resulting Sens is: 0.9895
The fit for each individual model on the Sens is:
  method      Sens      SensSD
svmPoly 0.9986372 0.0009085423
nnet     1.0000000 0.0000000000
C5.0     0.9858248 0.0043332664
nb       1.0000000 0.0000000000
> enstackpredictions <- predict(ensmodel2, testData, na.action = na.omit)
> cmBank2 <- confusionMatrix(ensstackpredictions, testData$y, mode="everything", positive = "no")
> cmBank2
Confusion Matrix and Statistics

              Reference
Prediction no yes
no      894   83
yes     23   29

              Accuracy : 0.897
              95% CI : (0.8768, 0.9149)
No Information Rate : 0.8912
P-Value [Acc > NIR] : 0.2941

              Kappa : 0.3057
McNemar's Test P-Value : 1.001e-08

              Sensitivity : 0.9749
              Specificity : 0.2589
Pos Pred Value : 0.9150
Neg Pred Value : 0.5577
Precision : 0.9150
Recall : 0.9749
F1 : 0.9440
Prevalence : 0.8912
Detection Rate : 0.8688
Detection Prevalence : 0.9495
```

```
> CostDTModel
```

```
Call:
C5.0.formula(formula = y ~ ., data = trainingData, cost = cmatrix)
```

```
Classification Tree
Number of samples: 3090
Number of predictors: 19
```

```
Tree size: 61
```

```
Non-standard options: attempt to group attributes
```

```
Cost Matrix:
```

```
      no yes
no     0  10
yes    1   0
```

```
Confusion Matrix and Statistics
```

```
              Reference
Prediction no yes
no      605  31
yes     312  81
```

Profit= $81 \cdot 10 - (312 + 81) \cdot 1 = 417$

Ensemble method can improve model accuracy by using bagging, booting, and stacking. In this case, I used resample, stack models to improve models. I find the profit increase. (no more than 425, I did not find reasons.)